

Aligned models cave to all four cue
types on questions they otherwise
answer correctly

